



Performance Playbook per Kernel, Loop-Level HIP

Speed-Up vs Numba

Qwen3.8-27B

GPT-OSS-120B



Tokens Spent

