



Performance Playbook per Kernel, Scientific Computing

Speed-Up vs Numba

Qwen3.8-27B

GPT-OSS-120B



Tokens Spent

