

Speed-up under kernel and repository framing, aggregated per kernel

○ kernel ● repo

GPT-OSS-120B / C

Qwen3.8 / C

n=2/2

n=2/2

2

3

geomean of each kernel's last submission (x, log scale)

