


llr40: best agent speed-up per kernel, C against Fortran (vs numba)



Values are UNVETTED: the implausible-speed-up check never fired. One graded aggregate per kernel and language, carrying no interval: the judge's repeat samples are not in this artifact, so SC15 rule 5 cannot be met per kernel here. 0 roster kernel(s) with no submission against this reference: none

• C (best over C arms) • Fortran (best over Fortran arms) ---- 1.0x (no change)