

v1l1w2: token cost per kernel

Median tokens per episode, with 95% bootstrap intervals. Log axis: spend spans decades.

