

Speedup over PyTorch (1 GPU)

Weak Scaling

Strong Scaling






MoE Router

SDPA

All-to-All

Vocab Softmax

Geomean



— Ideal Scaling of the PyTorch Baseline
— Qwen3.8-27B, No Skill Packet (n=15)
— GPT-OSS-120B, No Skill Packet (n=13)

GPUs P

— Kimi-K2.7-Code, No Skill Packet (n=17)
— PyTorch Distributed (n=20)