

Speedup over PyTorch (1 GPU)

Weak Scaling

Strong Scaling











MoE Router

SDPA

All-to-All

Vocab Softmax

Geomean

GPUs P

-- Ideal Scaling of the PyTorch Baseline

—○— Qwen3.8-27B, No Skill Packet (n=15)

—○— GPT-OSS-120B, No Skill Packet (n=13)

—○— Kimi-K2.7-Code, No Skill Packet (n=17)

—×— PyTorch Distributed (n=20)